

IP-Based Actions

IP-based actions are similar to BIND 9's RPZ-IP triggers and corresponding actions. They are similar to the existing local-zone or other tag-based actions, but defined for particular IP-netblocks. If the IP(v6/v4) address of an AAAA or A in the answer section matches one of the specified netblocks, the corresponding action will apply to the response. For example, it may change the address to a different one (redirect) or make the resolution result in NXDOMAIN.

This note describes details of proposed implementation of IP-based actions in Unbound.

Configuration Syntax

IP-based actions require a new (optional) Unbound module placed in the configuration (see the implementation note below). So the `module-c fig` option should be specified explicitly:

```
module-c fig "unbound-ip-action"
```

If DNSSEC-validation is proposed, it should be specified as follows:

```
module-c fig "unbound-ip-action"
```

We also introduce three configuration options for the `unbound-ip-action` module:

response-ip: <IP-netblock> <action>
If the IP address in an AAAA or A RR in the answer section matches the specified IP-netblock, the specified action will be applied. The semantics are the same as that for `action` (see below).

response-ip-data: <IP-netblock> <action> <data>
This specifies the action data for the specified IP-netblock. The "record string" is similar to that of `action` but it must be either AAAA, A, or CNAME. If the IP-netblock is an IPv6/IPv4 AAAA/A, respectively unless it's CNAME, which can be used for both netblocks). If it's CNAME, there must be more than one `response-ip-data` for the same IP-netblock. Also, CNAME and other types of records must not coexist for the same IP-netblock. The textual domain name for the CNAME does not have to be explicitly terminated with a dot (.); the origin is assumed to be the origin for the name.

response-ip-tag: <IP-netblock> <action> <tag>

Assign tags to response IP-netblocks. If the IP-netblock is AAAA, NOERR in the answer section is not needed. If the specified IP-netblock, then the specified tag is assigned to the response IP-netblock. If the tag is not specified, then the tag is assigned to the response IP-netblock. If the tag is not specified, then the tag is assigned to the response IP-netblock.

acce -c
acce -c
e e-i
defined for the

e e-i and e e-i can also be specified for a view, just like
l cal- e and l cal-

If multiple e e-i are specified for the same IP-netblock (in different lines), all but the one that appears first will be ignored (following the similar case for acce -c l- ag). If it's not a configuration error, it's quite unlikely that the result is the intended one and would be meaningless in practice.

The e e-i - options require the e-i module; if any of these options are specified but the e-i module is not loaded, the configuration will refuse to load the configuration. In fact, some part of the configuration (modifying cache answers) would work without the e-i module, but such an incomplete configuration would be even harmful. It should be better to refuse to load the configuration in such a case.

Exceptions on response

Actions specified for e e-i are different from those for local zones. In the former there is no point of having a tag as "e" because it is not used. Because of this difference, the following actions are not supported:

- a ic, ef e-i - options are not supported for e e-i - options.
- de is non-conditional, i.e., it always results in a response. It can be used for queries.

On the other hand, actions specified in a tag that has a matching tag with e e-i - ag are those that are "used" for response in listed above, since acce -c l- ag-a can be shared in local zones. In case of local zones, the behavior for response-ip-tags is not in case of local zones. If there is no corresponding e e-i -da a will generally result in NOERROR/NODATA instead of NXDOMAIN, since the response-ip data are inherited from the type specific, and non-existence of data doesn't indicate anything about the existence of the name itself. For example, if the matching tag action is static but there is no data for the corresponding response-ip

configuration, then the result will be NOERROR. NXDOMAIN is the only case where NXDOMAIN is returned is when an authoritative domain action applies.

Notes on Detailed Behavior

No AA Bit

If the original answer is tweaked due to a response-ip action, the final response sent to the client will never have the AA (authoritative answer) bit on in the header section. This is different from some cases of local-zone and access-control-tag actions (faked NXDOMAIN and local/redirect data), where the final response is considered “authoritative” and has the AA bit on. This behavior is consistent with BIND 9 RPZ in that case it’s not specific to RPZ-IP triggers).

Answer Records Only

IP-based actions apply only to records in the answer section. Even if the IP address of an AAAA or A record in the additional (or, unofficial practice, authority) section matches an IP-netblock of some response-ip options, those records (or the response itself) won’t be modified.

Multi-RR Case

If an AAAA or A RRset in the answer section of a query contains multiple addresses and one or more of the addresses are subject to response-ip processing, the action for the matching address in the answer message will be progressively used. For example, if an answer AAAA RRset contains the following IPv6 addresses in this order:

```
2001:db8::1
2001:db8::2
2001:db8::3
```

and the following IP-based actions are configured:

```
e      e-i : 2001:db8::1/128 authoritative domain
e      e-i : 2001:db8::2/128 redirect
e      e-i -data 2001:db8::bad
```

then the first matching action (always authoritative) will be used and the response will be converted to NXDOMAIN (with no answer records). If there are multiple matching redirect actions for different IP-netblocks of the same answer RRset, only the first matching redirect data will be used, and the resulting response will contain only one address (which is the matching redirect data). Since it appears that Unbound caches an RRset in the order it receives from the upstream server, the same action will be applied to subsequent responses to the same query when they are directly answered from the cache.

This behavior largely follows the behavior of BIND 9’s RPZ-IP triggers. Note, however, that the order of the answer RRs from remote servers is not always predictable, so it is also

unpredictable which `e` `e-i` action is used when there are multiple candidates. This is different from BIND 9 RPZ-IP where the triggers are applied to sorted RRs by default and the matching rule is generally preferred for the same.

CNAME Chasing

If a CNAME is specified as the target for an IP-based action, Unbound will automatically chase the CNAME target until it receives an answer (whether positive or negative, or some other error), and return the complete CNAME chain to the end client. For example, for the above original AAAA answer, we can define a redirect response-ip action as follows:

```
e      e-i : 2001:db8::1/128 edi ec
e      e-i -da a. 2001:db8::1/128 CNAME a ge .e am le.
```

then Unbound will resolve AAAA (either by retrieving it from the cache), and include the result in the final answer as follows. An example final answer is:

```
< igi al ame> 2001:db8::1/128 CNAME a ge .e am le.
a ge .e am le. 2001:db8::1/128 CNAME a ge .e am le.
```

Unlike “CNAME-based redirect response-ip” zone actions, CNAME chasing will take place even if the redirect data is used for a `control-tag-action` that is not redirect but can use local data (e.g. static). For example, in addition to the above configuration, if there’s a following tag setup:

```
acce -c l- 192.0.2.1/24 "m ag" a ic
e      e-i - ag: 2001:db8::1/128 "m ag"
```

then, if the client at 192.0.2.1 sends a query for the corresponding response-ip data, the target will be chased (in retrospect, this would have been more intuitive for the local zone case, too).

Even if the CNAME target RRset has a valid RRSIG, the RRSIG won’t be included in the final response to the client, nor the AD bit will be set in the response, regardless of whether or not the original query sets the DO or CD bit (see also “DNSSEC implications” below). This behavior is consistent with BIND 9.

Type-ANY query suppresses this chasing. For instance, in the above example configuration if the query type is ANY, the answer will only contain the CNAME RR. This behavior is consistent with BIND 9.

Note: Chasing CNAME targets for an IP-based action may be especially expensive in terms of performance (see implementation notes below). It’s probably advisable to avoid this configuration whenever possible.

No Recursive Application

This proposed implementation does not try to apply IP-based actions “recursively”; that is, it does not apply an IP-based action to the IP address specified as a result of data of a redirect

response-ip action. For example, we apply a redirect action for 2001:db8::3/128 used in the previous example and redirect to 2001:db8::bad. Then, even if there is another IP-based action which 2001:db8::3 matches, that action won't apply. (This is the same as BIND 9's RPZ-IP.)

The same restriction applies to the address of a CNAME when the CNAME is the data of an IP-based action. For example, the example of the CNAME chain shown as

```
< igi al ame> CNAME .e am le.  
a ge .e am le. AAAA 2001:db8::ffff
```

even if there is another IP-based action which 2001:db8::ffff matches, that action won't apply. In this case, only the CNAME will be returned to the client in order to avoid including an IP address in the answer (in this example, 2001:db8::ffff) that would otherwise be subject to an IP-based action. It will also help avoid having a weird corner cases like a CNAME chain (e.g., consider the case where 2001:db8::ffff is redirect to < igi al ame> a ge .e am le. Note that a sophisticated client (such as a local caching server using this Unbound as a "forwarder") could re-query for the CNAME target when it gets the incomplete CNAME chain. In this case the intended IP-based action will apply and that client will get an answer that the administrator of this action would probably envision. For traditional BIND resolver, an incomplete CNAME chain effectively means resolution failure. This is not ideal, but it would be acceptable in practice as in most cases the presence of such action would be to avoid returning a specific IP address.

In any case, such a situation is considered a local configuration error, and Unbound leaves an informational level of log messages when it detects the situation.

```
CNAME a ge ffffff e-i ac i ld be  
e e-i ac i ;
```

This is also consistent with BIND 9 RPZ, which never applies RPZ rules more than once. In this implementation we defer from it to keep the implementation simpler, but we may want to revisit it (this is also different from what's written in draft-vixie-dns-rpz).

Similar to the above cases, even if there is a local-zone action to which a ge .e am le would be subject, it won't apply. In this case the above AAAA RR will be included in the answer to the client (it's no different from how a redirected CNAME target in local-zone or tag-based actions works in general).

This is different from BIND 9 RPZ, which never applies RPZ rules more than once. In this implementation we defer from it to keep the implementation simpler, but we may want to revisit it (this is also different from what's written in draft-vixie-dns-rpz).

No Override for Other Local

Similar to the previous subsection, IP-based actions will not apply to local-zone or local-data or tag-based actions. For example, if we have the following response-ips configuration:

```
l cal- e: e am edi ec
l cal-da a: e a c m 192.0.2.1
e e-i : 192.0.2.0/24
```

then a query for example.com/A will be answered from the local-zone with IP 192.0.2.1. Even if it would match the response-ips configuration, it will not be overridden. This is consistent with BIND 9 RPZ.

Authority and Additional Sections

If an AAAA or A RR of the answer section of a response is modified due to an IP-based action, the authority and additional sections of the resulting final response sent to the client will be cleared (an EDNS OPT RR may still be added to the additional section). This is the case regardless of the action type, even if it's `edi ec` or `d mai` variant. This is different from BIND 9 RPZ: it adds an SOA of the corresponding RPZ to the authority section for negative answers and to the additional section for NXDOMAIN.

DNSSEC Implications

If an original response is modified due to an IP-based action, the resulting final response will never have the AD bit on even if the original response was DNSSEC-validated. Any RRSIG RRsets for the modified RRset will be removed from the answer section; however, if the original response is a CNAME chain and some of the CNAMEs have RRSIGs, these RRSIG will be kept in the final response (note that CNAMEs can never be subject to an IP-based action). This behavior is compatible with BIND 9 RPZ.

Multi-Level Netblock Matching

We will (eventually) try to implement multi-level matching: if the best (longest) matching IP-prefix does not have a matching tag for the client but a less-specific matching IP-netblock has a matching tag, the action for the less-specific IP-netblock should apply. For example, if we have the following two response-ips configurations:

```
e e-i - ag: 192.0.2.128/28 " ag1 ag2"
e e-i - ag: 192.0.2.255/32 " ag3"
```

and if an A RR in the answer section has IP 192.0.2.255 but tags for the client include "tag1" but not "tag3", then the action for 192.0.2.128/28 should apply even if the best matching netblock is 192.0.2.255/32.

This ideal behavior is consistent with how tags for local-zones match. But this may need non-trivial extensions to existing Unbound utilities, while such a multi-level setup is supposed to be quite rare. So we may skip this behavior and always consider the best matching netblock (if there is no matching tag, treat it as there is no matching netblock) in the initial implementation.

This is a **TODO** item as of this writing (as of February 10, 2017, this “ideal behavior” is not implemented).

Type ANY Query

The primarily intended use of local-zone actions is to apply them to type AAAA or A queries. But it should also work for type ANY queries. If an answer contains an AAAA or A record. This is consistent with BIND 9.10.1. In this case, records of the original answer will be removed except the one that triggered the action for which it may be replaced (in case of a redirect action) or a new record will be added. Types of actions including (but not limited to) variants). This behavior is consistent with BIND 9.10.1.

Inform Log

If an `if_mori` filter is applied to a message will be logged. The log message contains information about the filter and the action. The log message is similar to log messages of `if_mori` and `if_m_de` actions. The log message contains the message ID and the netblock instead of the domain name. An example log message is as follows:

```
if : 192.0.2.1/24 15005 if.f.e a.m l.e.c.t.u.r.e A
IN
```

Overview of Implementation Design

Unlike other local-zone and tag-based actions, whether to apply an IP-based action cannot be determined completely locally since it depends on the result of normal resolution. Also, since the “CNAME chasing” behavior will require additional sub-queries, we cannot just tweak the answer immediately before sending it to the client.

So we chose to implement the main functionality as a separate Unbound module named `response IP` (meaning “response IP”). It’s intended to be placed on top of the module stack (usually immediately above the iterator module) and works as a filter for resolution results. Its `resolve()` function first passes the state to the lower module to resolve the original query. When it’s done, it checks the answer section for AAAA or A records that may require IP-based action processing. If some action needs to be applied, the module modifies the answer accordingly (changing the IP address to a different one, converting the whole answer to NXDOMAIN, etc.). It usually completes the module. The exception is the case where a redirect to CNAME is triggered, which causes a sub-query to be triggered to resolve the CNAME target, and the original query state is held until the completion of the sub-query, the module’s `resolve()` function is called. It appends the resolved target records to the (already tweaked) answer section of the original answer. Finally, the control comes back to the `resolve()` function, which completes the module processing.

The `response IP` module will need to access part of `acl_add` of the client that triggers a particular query so that it can apply tag-based per-view actions. So we’ll need to extend the existing `me_h_e_client()` function so that it stores the sub part of `acl_add` and stores it in

the `module_acl` a new structure named `acl_client` is introduced for this purpose (we didn't like to reuse `acl_add` in the unbound daemon, as it was clearly not module-boundary friendly; otherwise we would have just used `acl_add` itself instead of a new structure). The module state comparison function also has to be updated so that it takes into account `acl_client`. Note that this can lead to more external queries for the same qname and more ACL entries are needed for many different clients.

We also need to update the code that answers queries directly from cache (mainly in `daemon/lookup.c`) if the answer to the query is already cached and it has to be modified by an IP-based action. In terms of observed behavior this could be avoided if we always call the module stack to get the answer; however, it's quite likely to be unacceptable in terms of performance since it would require additional function calls and many more data copies. (We won't be able to use the format cached data directly as the resolution needs to go to the generic `idea` module.)

Our implementation updates the `answer_from_cache()` function for this purpose. Before encoding the answer RRset, it now checks if an IP-based action should apply, and if so, tweaks the answer accordingly. Again, the tricky case is a redirection. In this case, it first needs to see if the CNAME target is cached, but to do so it will first make a local copy of the original RRsets, tweaks the answer, and releases acquired locks (it needs to release the lock for the additional cache lookup, and so it needs to make a copy of the cache data as the reference to the cache is not protected by the cache). Also, if the CNAME target is not cached, it will behave as if the original cache lookup failed to trigger the usual resolution (and processing by the `idea` module). These are complicated and inefficient, but seem to be unavoidable costs.